

# Business Understanding

The company is expanding into both privatized and commercial aviation to diversify its portfolio. They would like to understand which aircraft types have the lowest risk. The "AviationData.csv" dataset provided by the NTSB, that will be explored in this notebook, contains information on selected accidents from 1962 to 2023. These are incidents that have taken place within the United States, its territories and possessions, and in international waters.

## Key/focus questions:

- Which aircraft types have the lowest accident severity?
- Is a private or commercial flight riskier?
- Are certain conditions (e.g., weather, location) more associated with severe accidents?

**Audience:** Aviation division team and the company's stakeholders eg managers

```
In [1]: #importing the necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: #importing the dataset
df = pd.read_csv("C:/Users/PC/Desktop/School work/Projects/Phase1/Phase-1-Aviation-
C:\Users\PC\AppData\Local\Temp\ipykernel_13692\4218076357.py:2: DtypeWarning: Column
s (6,7,28) have mixed types. Specify dtype option on import or set low_memory=False.
df = pd.read_csv("C:/Users/PC/Desktop/School work/Projects/Phase1/Phase-1-Aviation
-Project/data/AviationData.csv", index_col = 0, encoding='cp1252')
```

# Understanding the Dataset

```
In [3]: #Looking into the datasets shape
df.shape
```

```
Out[3]: (88889, 30)
```

```
In [4]: #information on the dataset
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
Index: 88889 entries, 20001218X45444 to 20221230106513
Data columns (total 30 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Investigation.Type                    88889 non-null  object
1   Accident.Number                      88889 non-null  object
2   Event.Date                          88889 non-null  object
3   Location                            88837 non-null  object
4   Country                             88663 non-null  object
5   Latitude                           34382 non-null  object
6   Longitude                           34373 non-null  object
7   Airport.Code                        50132 non-null  object
8   Airport.Name                        52704 non-null  object
9   Injury.Severity                     87889 non-null  object
10  Aircraft.damage                      85695 non-null  object
11  Aircraft.Category                   32287 non-null  object
12  Registration.Number                 87507 non-null  object
13  Make                               88826 non-null  object
14  Model                              88797 non-null  object
15  Amateur.Built                      88787 non-null  object
16  Number.of.Engines                  82805 non-null  float64
17  Engine.Type                        81793 non-null  object
18  FAR.Description                    32023 non-null  object
19  Schedule                          12582 non-null  object
20  Purpose.of.flight                  82697 non-null  object
21  Air.carrier                        16648 non-null  object
22  Total.Fatal.Injuries                77488 non-null  float64
23  Total.Serious.Injuries              76379 non-null  float64
24  Total.Minor.Injuries                76956 non-null  float64
25  Total.Uninjured                     82977 non-null  float64
26  Weather.Condition                  84397 non-null  object
27  Broad.phase.of.flight               61724 non-null  object
28  Report.Status                       82505 non-null  object
29  Publication.Date                    75118 non-null  object
dtypes: float64(5), object(25)
memory usage: 21.0+ MB

```

```

In [5]: #data types of the columns
        df.dtypes

```

```
Out[5]: Investigation.Type      object
        Accident.Number      object
        Event.Date           object
        Location              object
        Country               object
        Latitude              object
        Longitude             object
        Airport.Code          object
        Airport.Name          object
        Injury.Severity       object
        Aircraft.damage       object
        Aircraft.Category     object
        Registration.Number   object
        Make                  object
        Model                 object
        Amateur.Built         object
        Number.of.Engines     float64
        Engine.Type           object
        FAR.Description       object
        Schedule              object
        Purpose.of.flight     object
        Air.carrier           object
        Total.Fatal.Injuries  float64
        Total.Serious.Injuries float64
        Total.Minor.Injuries  float64
        Total.Uninjured       float64
        Weather.Condition     object
        Broad.phase.of.flight object
        Report.Status         object
        Publication.Date      object
        dtype: object
```

```
In [6]: #Looking into the first 5 rows of the dataset to get an idea of the data
        df.head()
```

Out[6]:

Event.Id	Investigation.Type	Accident.Number	Event.Date	Location	Country
20001218X45444	Accident	SEA87LA080	1948-10-24	MOOSE CREEK, ID	United States
20001218X45447	Accident	LAX94LA336	1962-07-19	BRIDGEPORT, CA	United States
20061025X01555	Accident	NYC07LA005	1974-08-30	Saltville, VA	United States
20001218X45448	Accident	LAX96LA321	1977-06-19	EUREKA, CA	United States
20041105X01764	Accident	CHI79FA064	1979-08-02	Canton, OH	United States

5 rows × 30 columns

In [7]: df.describe()

Out[7]:

	Number.ofEngines	Total.Fatal.Injuries	Total.Serious.Injuries	Total.Minor.Injuries	Tot
count	82805.000000	77488.000000	76379.000000	76956.000000	8
mean	1.146585	0.647855	0.279881	0.357061	
std	0.446510	5.485960	1.544084	2.235625	
min	0.000000	0.000000	0.000000	0.000000	
25%	1.000000	0.000000	0.000000	0.000000	
50%	1.000000	0.000000	0.000000	0.000000	
75%	1.000000	0.000000	0.000000	0.000000	
max	8.000000	349.000000	161.000000	380.000000	

# Cleaning the dataset

Using the head function and looking at the different columns of the dataset, some columns are not required moving forward.

**Why?** This is because they are not part/helpful in the analysis and will not interfere when forming a data driven conclusion. The columns such as Event Id and Accident number are just for administration purposes and Report status and publication date would be desired for a different type of analysis

In [8]: *#Getting the column names in the dataset*

```
df.columns
```

Out[8]: Index(['Investigation.Type', 'Accident.Number', 'Event.Date', 'Location',  
'Country', 'Latitude', 'Longitude', 'Airport.Code', 'Airport.Name',  
'Injury.Severity', 'Aircraft.damage', 'Aircraft.Category',  
'Registration.Number', 'Make', 'Model', 'Amateur.Built',  
'Number.ofEngines', 'Engine.Type', 'FAR.Description', 'Schedule',  
'Purpose.of.flight', 'Air.carrier', 'Total.Fatal.Injuries',  
'Total.Serious.Injuries', 'Total.Minor.Injuries', 'Total.Uninjured',  
'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',  
'Publication.Date'],  
dtype='object')

In [9]: *#forming a list of the unrequired columns*

```
unwanted_columns = ['Latitude', 'Longitude', 'Airport.Code',  
                    'Airport.Name', 'Registration.Number', 'Accident.Number', 'Report.Status',  
                    'Publication.Date']  
df = df.drop(columns = unwanted_columns)
```

In [10]: df.head()

Out[10]:

	Investigation.Type	Event.Date	Location	Country	Injury.Severity	Ai
	Event.Id					
	20001218X45444	Accident	1948-10-24	MOOSE CREEK, ID	United States	Fatal(2)
	20001218X45447	Accident	1962-07-19	BRIDGEPORT, CA	United States	Fatal(4)
	20061025X01555	Accident	1974-08-30	Saltville, VA	United States	Fatal(3)
	20001218X45448	Accident	1977-06-19	EUREKA, CA	United States	Fatal(2)
	20041105X01764	Accident	1979-08-02	Canton, OH	United States	Fatal(1)

5 rows × 22 columns



In [11]: *# checking for columns with missing values*

```
missing_count = df.isna().sum()  
missing_count
```

```
Out[11]: Investigation.Type      0
          Event.Date            0
          Location              52
          Country              226
          Injury.Severity      1000
          Aircraft.damage      3194
          Aircraft.Category    56602
          Make                 63
          Model                92
          Amateur.Built        102
          Number.of.Engines    6084
          Engine.Type          7096
          FAR.Description      56866
          Schedule            76307
          Purpose.of.flight    6192
          Air.carrier          72241
          Total.Fatal.Injuries 11401
          Total.Serious.Injuries 12510
          Total.Minor.Injuries 11933
          Total.Uninjured      5912
          Weather.Condition    4492
          Broad.phase.of.flight 27165
          dtype: int64
```

```
In [12]: #checking for the percentages of columns with missing values
#if it is greater than 50%, the columns are to be dropped

row_number = len(df)

missing_percentage = (missing_count / row_number) * 100
missing_percentage.sort_values(ascending = True)
```

```
Out[12]: Investigation.Type      0.000000
          Event.Date            0.000000
          Location              0.058500
          Make                 0.070875
          Model                0.103500
          Amateur.Built        0.114750
          Country              0.254250
          Injury.Severity      1.124999
          Aircraft.damage      3.593246
          Weather.Condition    5.053494
          Total.Uninjured      6.650992
          Number.of.Engines    6.844491
          Purpose.of.flight    6.965991
          Engine.Type          7.982990
          Total.Fatal.Injuries 12.826109
          Total.Minor.Injuries 13.424608
          Total.Serious.Injuries 14.073732
          Broad.phase.of.flight 30.560587
          Aircraft.Category    63.677170
          FAR.Description      63.974170
          Air.carrier          81.271023
          Schedule            85.845268
          dtype: float64
```

In [13]: *#dropping the columns with high percentage of missing values*

```
df = df.drop(columns = ['Schedule', 'Air.carrier', 'FAR.Description'])
df.head()
```

Out[13]:

	Investigation.Type	Event.Date	Location	Country	Injury.Severity	Airline

Event.Id

20001218X45444	Accident	1948-10-24	MOOSE CREEK, ID	United States	Fatal(2)	
20001218X45447	Accident	1962-07-19	BRIDGEPORT, CA	United States	Fatal(4)	
20061025X01555	Accident	1974-08-30	Saltville, VA	United States	Fatal(3)	
20001218X45448	Accident	1977-06-19	EUREKA, CA	United States	Fatal(2)	
20041105X01764	Accident	1979-08-02	Canton, OH	United States	Fatal(1)	



In [14]: df.shape

Out[14]: (88889, 19)

In [15]: *#Checking for duplicates in the data*

```
df.duplicated().value_counts()
```

Out[15]:

```
False    88854
True       35
Name: count, dtype: int64
```

In [16]: *#dropping the duplicates*

```
df = df.drop_duplicates()
df.shape
```

Out[16]: (88854, 19)

In [17]: *#rechecking duplicates*

```
df.duplicated().value_counts()
```

Out[17]:

```
False    88854
Name: count, dtype: int64
```

## Handling Missing Values

Working according to the data types i.e, categorical or numerical, and the outliers.

In [18]: `df.dtypes`

```
Out[18]: Investigation.Type      object
Event.Date                    object
Location                      object
Country                       object
Injury.Severity               object
Aircraft.damage               object
Aircraft.Category             object
Make                          object
Model                         object
Amateur.Built                 object
Number.of.Engines             float64
Engine.Type                   object
Purpose.of.flight             object
Total.Fatal.Injuries           float64
Total.Serious.Injuries         float64
Total.Minor.Injuries           float64
Total.Uninjured               float64
Weather.Condition              object
Broad.phase.of.flight          object
dtype: object
```

In [19]: `df.isna().sum().sort_values(ascending = True)`

```
Out[19]: Investigation.Type      0
Event.Date                    0
Location                      52
Make                          63
Model                         92
Amateur.Built                 102
Country                       226
Injury.Severity               998
Aircraft.damage               3191
Weather.Condition              4490
Total.Uninjured               5908
Number.of.Engines             6080
Purpose.of.flight             6189
Engine.Type                   7093
Total.Fatal.Injuries           11398
Total.Minor.Injuries           11926
Total.Serious.Injuries         12503
Broad.phase.of.flight          27158
Aircraft.Category             56577
dtype: int64
```

## Numerical columns

- First step will be describing the numerical columns and see if their mean and any outliers
- Fill the missing values using an appropriate and logical manner

After handling the missing data:



- Plotting graphs to visualise these outliers
- Handling the outliers if present
- Plot a graph to indicate the difference

In [20]: `df.columns`

Out[20]: Index(['Investigation.Type', 'Event.Date', 'Location', 'Country',  
'Injury.Severity', 'Aircraft.damage', 'Aircraft.Category', 'Make',  
'Model', 'Amateur.Built', 'Number.of.Engines', 'Engine.Type',  
'Purpose.of.flight', 'Total.Fatal.Injuries', 'Total.Serious.Injuries',  
'Total.Minor.Injuries', 'Total.Uninjured', 'Weather.Condition',  
'Broad.phase.of.flight'],  
dtype='object')

In [21]: *#creating a list of the columns which are float as indicated above*  
float\_columns = ['Number.of.Engines', 'Total.Fatal.Injuries', 'Total.Serious.Injuri  
float\_columns

Out[21]: ['Number.of.Engines',  
'Total.Fatal.Injuries',  
'Total.Serious.Injuries',  
'Total.Minor.Injuries',  
'Total.Uninjured']

- Due to an error faced later as your filling in the missing values, converting to numeric is required since the columns are float
- This means using the `pd.to_numeric` function

In [22]: `for col in float_columns:  
df[col] = pd.to_numeric(df[col], errors='coerce')`

In [23]: `for col in float_columns:  
print(df[col].dtype)`

float64  
float64  
float64  
float64  
float64

In [24]: *#use a for loop to describe each of the columns in the above list to see their prop*  
  
`for col in float_columns:  
print(f"{col} description:")  
print(df[col].describe())  
print("\n") #having line between each description`

Number.of.Engines description:

count	82774.000000
mean	1.146556
std	0.446518
min	0.000000
25%	1.000000
50%	1.000000
75%	1.000000
max	8.000000

Name: Number.of.Engines, dtype: float64

Total.Fatal.Injuries description:

count	77456.000000
mean	0.647826
std	5.487038
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	349.000000

Name: Total.Fatal.Injuries, dtype: float64

Total.Serious.Injuries description:

count	76351.000000
mean	0.279892
std	1.544285
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	161.000000

Name: Total.Serious.Injuries, dtype: float64

Total.Minor.Injuries description:

count	76928.000000
mean	0.357061
std	2.235891
min	0.000000
25%	0.000000
50%	0.000000
75%	0.000000
max	380.000000

Name: Total.Minor.Injuries, dtype: float64

Total.Uninjured description:

count	82946.000000
mean	5.318159
std	27.891441
min	0.000000
25%	0.000000
50%	1.000000
75%	2.000000

```
max          699.000000
Name: Total.Uninjured, dtype: float64
```

In [25]: *#calculating the median of the float columns*

```
for col in float_columns:
    print(f"Median for {col}:")
    print(df[col].median())
    print("\n")
```

```
Median for Number.of.Engines:
1.0
```

```
Median for Total.Fatal.Injuries:
0.0
```

```
Median for Total.Serious.Injuries:
0.0
```

```
Median for Total.Minor.Injuries:
0.0
```

```
Median for Total.Uninjured:
1.0
```

From the results above:

- The number of engines is mostly 1. This is the suitable number to use to fill in the missing values in the rows for the Number of engine column. The outliers being 8 for example are for aircrafts which are used in the military. In this case, we are analysing for commercial or private aircrafts hence not considered.
- For injury based columns (the other 4 float columns), missing values means that the injuries were not collected/noted down. Therefore cannot fill in the missing values using the median/mean as it will not be accurate/correct assumption.

In [26]: *#filling in the missing values*

```
#filling in the no.of.engines column
df['Number.of.Engines'] = df['Number.of.Engines'].fillna(df['Number.of.Engines'].me
```

In [27]: *# other 4 columns that are to be filled with 0*

```
four_columns = ['Total.Fatal.Injuries', 'Total.Serious.Injuries', 'Total.Minor.Inju
for col in four_columns:
    df[col] = df[col].fillna(0)
```


```
In [28]: #checking if they have been filled  
df[float_columns].isnull().sum()
```

```
Out[28]: Number.of.Engines      0  
Total.Fatal.Injuries          0  
Total.Serious.Injuries        0  
Total.Minor.Injuries           0  
Total.Uninjured               0  
dtype: int64
```


- Able to handle the missing values in the numerical columns and now proceed to handle the outliers based on the results from the describe function

```
In [29]: # create a boxplot to visualise the outliers in each of the float columns in the li  
for col in float_columns:  
    fig, ax = plt.subplots(figsize=(8,4))  
    sns.boxplot(x=df[col], ax=ax)  
    ax.set_title(f'Boxplot for {col} (Before Handling the Outliers)')  
    plt.show()
```


Boxplot for Number.of.Engines (Before Handling the Outliers)



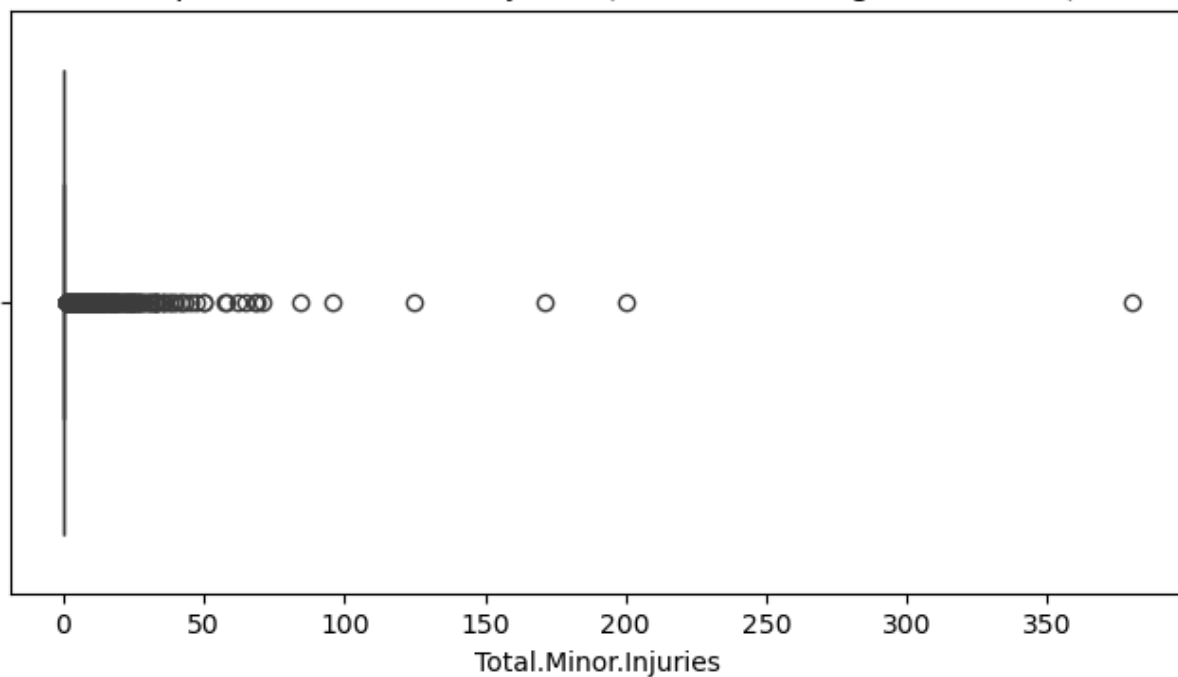
Boxplot for Total.Fatal.Injuries (Before Handling the Outliers)




Boxplot for Total.Serious.Injuries (Before Handling the Outliers)



Boxplot for Total.Minor.Injuries (Before Handling the Outliers)



Boxplot for Total.Uninjured (Before Handling the Outliers)



- We can now visually notice that there are outliers present in the numerical columns that need to be dealt with

In [30]: *# removing the outliers in the float columns*

```
for col in float_columns:  
    Q1 = df[col].quantile(0.25)  
    Q3 = df[col].quantile(0.75)  
    IQR = Q3 - Q1
```

```
# for correcting the dataframe without outliers present
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

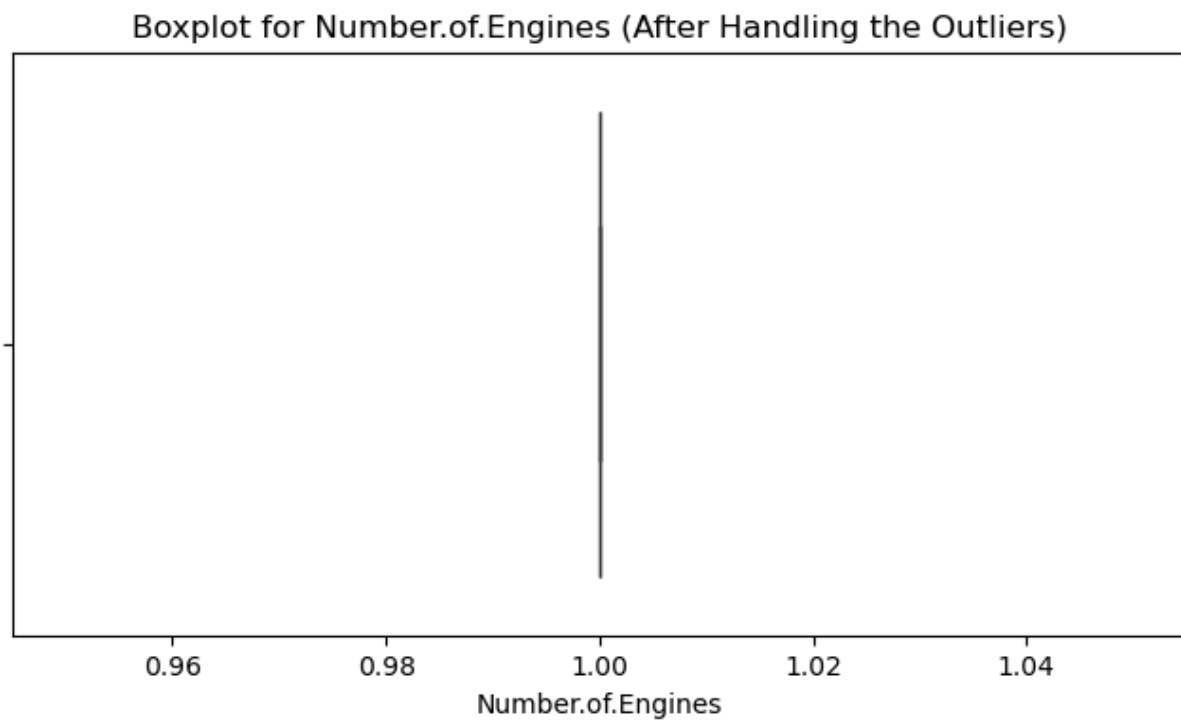
df = df[(df[col] >= lower_bound) & (df[col] <= upper_bound)]
```

In [31]: `df.shape`


Out[31]: (37805, 19)

In [32]: `# plotting after handling the columns outliers`


```
for col in float_columns:
    fig, ax = plt.subplots(figsize=(8,4))
    sns.boxplot(x=df[col], ax=ax)
    ax.set_title(f'Boxplot for {col} (After Handling the Outliers)')
    plt.show()
```



Boxplot for Total.Fatal.Injuries (After Handling the Outliers)




Boxplot for Total.Serious.Injuries (After Handling the Outliers)






Boxplot for Total.Minor.Injuries (After Handling the Outliers)



Boxplot for Total.Uninjured (After Handling the Outliers)



```
In [33]: # Lastly can describe and see if the outliers are still present
```

```
for col in float_columns:
    print(f"{col} description:")
    print(df[col].describe())
    print("\n") #having line between each description
```

Number.of.Engines description:

count	37805.0
mean	1.0
std	0.0
min	1.0
25%	1.0
50%	1.0
75%	1.0
max	1.0

Name: Number.of.Engines, dtype: float64

Total.Fatal.Injuries description:

count	37805.0
mean	0.0
std	0.0
min	0.0
25%	0.0
50%	0.0
75%	0.0
max	0.0

Name: Total.Fatal.Injuries, dtype: float64

Total.Serious.Injuries description:

count	37805.0
mean	0.0
std	0.0
min	0.0
25%	0.0
50%	0.0
75%	0.0
max	0.0

Name: Total.Serious.Injuries, dtype: float64

Total.Minor.Injuries description:

count	37805.0
mean	0.0
std	0.0
min	0.0
25%	0.0
50%	0.0
75%	0.0
max	0.0

Name: Total.Minor.Injuries, dtype: float64

Total.Uninjured description:

count	37805.000000
mean	1.480122
std	0.677462
min	0.000000
25%	1.000000
50%	1.000000
75%	2.000000

```
max          3.000000
Name: Total.Uninjured, dtype: float64
```

## Categorical data

```
In [34]: # checking the columns again
```


```
df.isna().sum().sort_values(ascending = True)
```

```
Out[34]: Investigation.Type          0
Event.Date                        0
Total.Uninjured                   0
Total.Minor.Injuries              0
Total.Serious.Injuries            0
Total.Fatal.Injuries              0
Number.of.Engines                 0
Location                          16
Make                              31
Model                             39
Country                           50
Amateur.Built                     54
Aircraft.damage                   703
Injury.Severity                   765
Weather.Condition                 1515
Purpose.of.flight                 1552
Engine.Type                       2188
Broad.phase.of.flight             11372
Aircraft.Category                 23196
dtype: int64
```

```
In [35]: # we can create a graph to visualise this
```

```
missing_counts = df.isna().sum()
missing_counts = missing_counts[missing_counts > 0].sort_values(ascending = True)

# Plot
plt.figure(figsize = (12,6))
sns.barplot(x = missing_counts.index, y = missing_counts.values)
plt.title('Missing Values Amount per Column')
plt.ylabel('Number of Missing Values')
plt.xlabel('Column Names')
plt.xticks(rotation = 90)
plt.show()
```



From the results above:

- The numerical columns with missing values have been handled and the categorical columns are yet to be handled
- The aircraft category column has a very high value of missing data, there needs to be dropped as it will not be helpful

```
In [36]: # we first drop the aircraft category column
df = df.drop(columns=['Aircraft.Category'])
df.head(2)
```

Out[36]:

	Investigation.Type	Event.Date	Location	Country	Injury.Severity	Aircr
Event.Id						
20020909X01562	Accident	1982-01-01	PULLMAN, WA	United States	Non-Fatal	
20020909X01559	Accident	1982-01-01	HOBBS, NM	United States	Non-Fatal	

```
In [37]: # after dropping the column, proceed to create a List of the columns that are object
object_columns = list(df.select_dtypes(include='object'))
object_columns
```

```
Out[37]: ['Investigation.Type',  
         'Event.Date',  
         'Location',  
         'Country',  
         'Injury.Severity',  
         'Aircraft.damage',  
         'Make',  
         'Model',  
         'Amateur.Built',  
         'Engine.Type',  
         'Purpose.of.flight',  
         'Weather.Condition',  
         'Broad.phase.of.flight']
```

```
In [38]: # from the list above get the columns which are missing values
```

```
missing_columns = []  
  
for col in object_columns:  
    if df[col].isna().any():  
        missing_columns.append(col)  
  
missing_columns
```

```
Out[38]: ['Location',  
         'Country',  
         'Injury.Severity',  
         'Aircraft.damage',  
         'Make',  
         'Model',  
         'Amateur.Built',  
         'Engine.Type',  
         'Purpose.of.flight',  
         'Weather.Condition',  
         'Broad.phase.of.flight']
```

```
In [39]: #now getting their unique values in the list above
```

```
for col in missing_columns:  
    print(f"Unique values in {col}:")  
    print(df[col].unique())  
    print("\n")
```

Unique values in Location:

```
['PULLMAN, WA' 'HOBBS, NM' 'TUSKEGEE, AL' ... 'Kildare, '
 'Rancho Cordova, CA' 'San Manual, AZ']
```

Unique values in Country:

```
['United States' 'GULF OF MEXICO' 'Puerto Rico' nan 'Bahamas'
 'Netherlands Antilles' 'CARIBBEAN SEA' 'Philippines' 'Mexico'
 'ATLANTIC OCEAN' 'Northern Mariana Islands' 'Turks And Caicos Islands'
 'Japan' 'Panama' 'PACIFIC OCEAN' 'American Samoa' 'Germany' 'Sudan'
 'United Kingdom' 'Venezuela' 'Canada' 'Switzerland' 'Argentina'
 'West Indies' 'Brazil' 'Trinidad And Tobago' 'Colombia' 'Fiji'
 'Netherlands' 'Ireland' 'Peru' 'Central African Republic'
 'British Virgin Islands' 'China' 'Saudi Arabia' 'Mauritius' 'Thailand'
 'Belgium' 'Uruguay' 'Guatemala' 'France' 'Taiwan' 'Iceland' 'Kazakhstan'
 'Ethiopia' 'Australia' 'Egypt' 'India' 'Spain' 'Greece'
 'Korea, Republic Of' 'Ecuador' 'Mozambique' 'Italy' 'Singapore'
 'Indonesia' 'Portugal' 'Finland' 'Guyana' 'Turkey' 'New Zealand'
 'Zimbabwe' 'Costa Rica' 'Federated States Of Micronesia' 'Norway' 'Gabon'
 'South Africa' 'Angola' 'Kenya' 'Guadeloupe' 'Ivory Coast'
 'Dominican Republic' 'Vietnam' 'Nepal' 'Sweden' 'Nigeria' 'Malaysia'
 'Cuba' 'Austria' 'Namibia' 'Czech Republic' 'Martinique' 'Denmark'
 'Chile' 'Belize' 'Bolivia' 'Reunion' 'United Arab Emirates' 'Ukraine'
 'Jordan' 'Nicaragua' 'Pakistan' 'Qatar' 'Poland' 'Israel' 'Jamaica'
 'Bangladesh' 'Hungary' 'Sint Maarten' 'Central Africa' 'Seychelles'
 'Rwanda' 'Afghanistan' 'Russia' 'South Korea' 'Iran' 'Lithuania'
 'Barbados' 'Eswatini' 'Antigua and Barbuda' 'Maldives' 'AY' 'Latvia'
 'Ghana' 'Macao' 'Luxembourg' 'Tanzania' 'Senegal' 'Papua New Guinea'
 'Cayman Islands' 'Solomon Islands' 'Mali' 'Turks And Caicos' 'Slovenia'
 'Cameroon' 'Bahrain' 'Nauru' 'Niue' 'UN' 'Antarctica' 'Bulgaria'
 'Morocco' 'Hong Kong' 'Romania' 'Suriname' 'Saint Barthelemy' 'Somalia'
 'Honduras' 'Guinea' 'Greenland' 'Estonia' 'Kyrgyzstan' 'Albania'
 'Croatia' 'Malta' 'South Sudan' 'Virgin Islands' 'Wallis and Futuna'
 'Lebanon' 'Saint Pierre and Miquelon' 'Georgia' "Côte d'Ivoire"
 'French Polynesia' 'Serbia' 'MU' 'Great Britain']
```

Unique values in Injury.Severity:

```
['Non-Fatal' 'Incident' 'Unavailable' nan 'Fatal' 'Serious' 'Minor']
```

Unique values in Aircraft.damage:

```
['Substantial' 'Destroyed' 'Minor' nan 'Unknown']
```

Unique values in Make:

```
['Cessna' 'Piper' 'Beech' ... 'CHILDS MICHAEL A' 'RUTAN'
 'JAMES R DERNOVSEK']
```

Unique values in Model:

```
['140' 'PA-28-161' 'V35B' ... 'KITFOX S5' '441' 'A330-243']
```

Unique values in Amateur.Built:

```
['No' 'Yes' nan]
```

Unique values in Engine.Type:

```
['Reciprocating' 'Turbo Shaft' 'Turbo Prop' 'Unknown' 'Turbo Jet'
 'Turbo Fan' nan 'Electric' 'Geared Turbofan']
```

Unique values in Purpose.of.flight:

```
['Personal' 'Instructional' 'Unknown' 'Ferry' 'Business'
 'Aerial Application' 'Executive/corporate' 'Aerial Observation'
 'Public Aircraft' 'Other Work Use' 'Positioning' nan 'Skydiving'
 'Flight Test' 'Air Race/show' 'Air Drop' 'Glider Tow'
 'Public Aircraft - Local' 'External Load' 'Banner Tow'
 'Public Aircraft - Federal' 'Public Aircraft - State' 'Firefighting'
 'Air Race show' 'PUBS' 'ASHO' 'PUBL']
```

Unique values in Weather.Condition:

```
['VMC' 'IMC' 'UNK' nan 'Unk']
```

Unique values in Broad.phase.of.flight:

```
['Takeoff' 'Approach' 'Landing' 'Taxi' 'Cruise' 'Climb' 'Descent'
 'Go-around' 'Maneuvering' 'Standing' 'Other' 'Unknown' nan]
```

In [40]: *#getting the mode of the unique values in the missing columns data*

```
for col in missing_columns:
    print(f"Mode for {col}: {df[col].mode()[0]}")
    print("\n")
```

Mode for Location: ANCHORAGE, AK

Mode for Country: United States

Mode for Injury.Severity: Non-Fatal

Mode for Aircraft.damage: Substantial

Mode for Make: Cessna

Mode for Model: 152

Mode for Amateur.Built: No

Mode for Engine.Type: Reciprocating

Mode for Purpose.of.flight: Personal

Mode for Weather.Condition: VMC

Mode for Broad.phase.of.flight: Landing

From the results above:

- We notice that the amateur built column has 2 options which is boolean (yes or no) therefore replace missing values with the mode
- The other columns have multiple options, hence will not be suitable to replace with the mode but rather "missing"

```
In [41]: #handling the missing data for the missing columns
```

```
for col in missing_columns:  
    if col == 'Amateur.Built':  
        df[col] = df[col].fillna(df[col].mode()[0])  
    else:  
        df[col] = df[col].fillna('Missing')
```

```
In [42]: df.isna().any()
```



```
Out[42]: Investigation.Type      False
         Event.Date            False
         Location              False
         Country               False
         Injury.Severity       False
         Aircraft.damage       False
         Make                  False
         Model                 False
         Amateur.Built         False
         Number.of.Engines     False
         Engine.Type           False
         Purpose.of.flight     False
         Total.Fatal.Injuries  False
         Total.Serious.Injuries False
         Total.Minor.Injuries  False
         Total.Uninjured       False
         Weather.Condition     False
         Broad.phase.of.flight False
         dtype: bool
```

```
In [43]: # ensuring the columns which are object are the same with the all the letters upper
         #avoiding Boeing and BOEING
         # Loop through all object columns and clean
         for col in df.select_dtypes(include='object').columns:
             df[col] = df[col].str.strip().str.upper()
```

```
In [44]: #converting the dataset that's been cleaned to csv to use for tableau

         df.to_csv('cleaned_aviation_data.csv', index=False)
```

## Data Analysis

- In this section the dataset will be narrowed down to the top 10 for filtering purposes
- This is to guide into drawing a well informed conclusion on which airplanes and approach is best for the company
- There will be use of diagrams such as bar graphs to form a visual conclusion
- Functions such as groupby will be used to combine/filter the dataset

```
In [45]: df.columns
```

```
Out[45]: Index(['Investigation.Type', 'Event.Date', 'Location', 'Country',
               'Injury.Severity', 'Aircraft.damage', 'Make', 'Model', 'Amateur.Built',
               'Number.of.Engines', 'Engine.Type', 'Purpose.of.flight',
               'Total.Fatal.Injuries', 'Total.Serious.Injuries',
               'Total.Minor.Injuries', 'Total.Uninjured', 'Weather.Condition',
               'Broad.phase.of.flight'],
              dtype='object')
```

```
In [46]: df.head(2)
```

Out[46]:

	Investigation.Type	Event.Date	Location	Country	Injury.Severity	Aircr
Event.Id						
20020909X01562	ACCIDENT	1982-01-01	PULLMAN, WA	UNITED STATES	NON-FATAL	SI
20020909X01559	ACCIDENT	1982-01-01	HOBBS, NM	UNITED STATES	NON-FATAL	SI

```
In [47]: # number of makes of aircrafts
number_of_makes = df['Make'].nunique()
print(f"Number of unique aircraft makes: {number_of_makes}")
```

Number of unique aircraft makes: 3311

```
In [48]: # number of models of aircrafts
number_of_models = df['Model'].nunique()
print(f"Number of unique aircraft models: {number_of_models}")
```

Number of unique aircraft models: 5365


## Frequency of accident over the years

- This will show and help us understand the trend from 1962-2023

```
In [49]: # the column Event Date needs to be converted to datetime
df['Event.Date'] = pd.to_datetime(df['Event.Date'], errors='coerce')
# need to create a column with just the year on its own which would be 1962 for exa
df['Year'] = df['Event.Date'].dt.year
```

```
In [50]: # for the plot, it needs to be grouped by year
accidents_per_year = df.groupby('Year').size()

plt.figure(figsize = (12,6))
sns.lineplot(x = accidents_per_year.index, y = accidents_per_year.values)
plt.title('Accident Frequency Over Years')
plt.xlabel('Year')
plt.ylabel('Number of Accidents')
plt.show()
```



From the line graph above you are able to notice that over the years, the number of accidents have decreased. This is possibly due to the improvements in technology, communication systems and safety protocols over the years. For the company, it would be an indication that investing in aviation is not a bad idea.

### Top 10 Makes of aircrafts involved in accidents

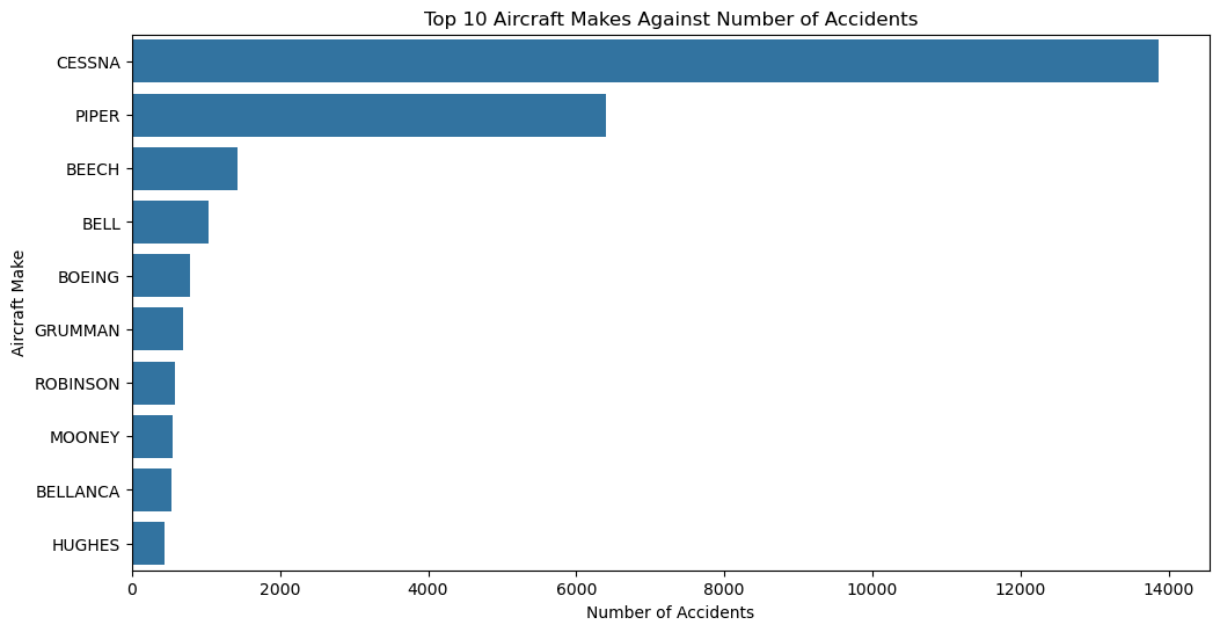
- This will help understand which makes to avoid

```
In [51]: top_makes = df['Make'].value_counts().head(10)
top_makes
```

```
Out[51]: Make
CESSNA      13858
PIPER       6406
BEECH       1422
BELL        1030
BOEING       779
GRUMMAN      692
ROBINSON     581
MOONEY       548
BELLANCA     531
HUGHES       441
Name: count, dtype: int64
```

```
In [52]: #plotting the graph for visualising

plt.figure(figsize = (12,6))
sns.barplot(x = top_makes.values, y = top_makes.index)
plt.title('Top 10 Aircraft Makes Against Number of Accidents')
plt.xlabel('Number of Accidents')
plt.ylabel('Aircraft Make')
plt.show()
```



From the horizontal graph above, it would make most sense to especially avoid Cessna with the highest followed by Piper. These makes makes (the top 2 mentioned) have averagely had the largest number of accidents which would not be a good investment for the company. It would put them and clientel at risk.

### Top 10 Models against the number of accident


- Helps us understand the models to avoid

```
In [53]: top_models = df['Model'].value_counts().head(10)
top_models
```

```
Out[53]: Model
152          1669
172          1051
172N           635
150           508
180           448
172M           410
PA-18          403
172P           396
PA-18-150      391
PA-28-140      383
Name: count, dtype: int64
```

```
In [54]: #plotting the graph for visualising

plt.figure(figsize = (12,6))
sns.barplot(x = top_models.values, y = top_models.index)
plt.title('Top 10 Aircraft Models Against Number of Accidents')
plt.xlabel('Number of Accidents')
plt.ylabel('Aircraft Model')
plt.show()
```



The 152, 172 and 172N models would be the riskiest and worst options to invest in as a company. This is because they hold the highest number of accidents. The last 5 models hold averagely the same amount of accidents; 172M - PA-28-140.

### Top 10 Total Uninjured by Aircraft Make

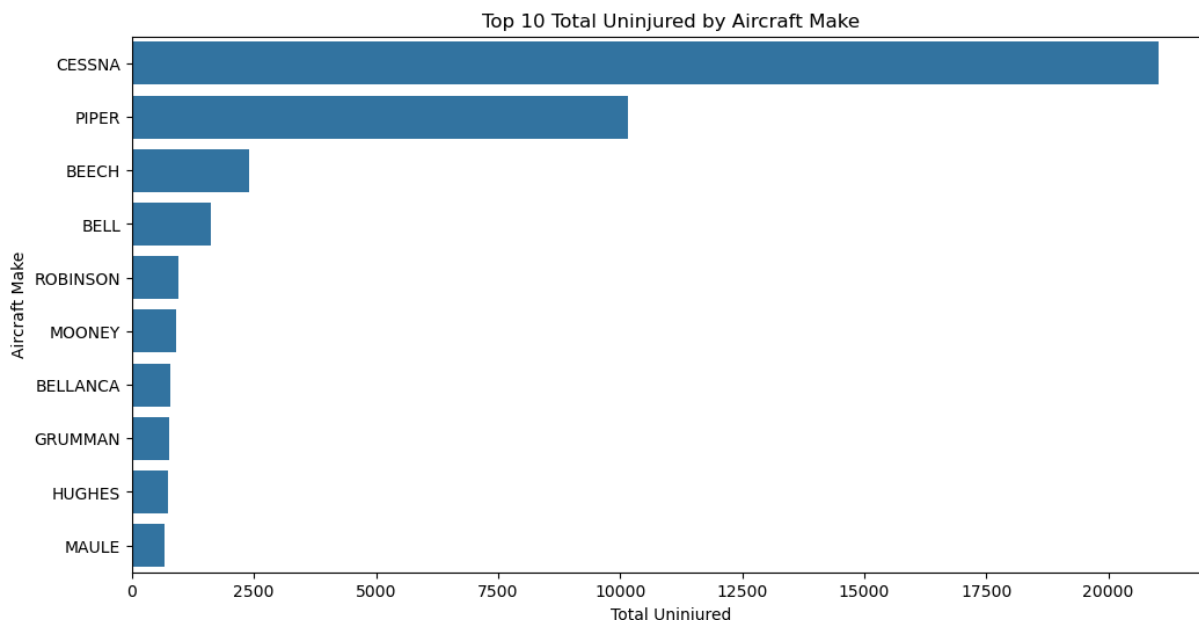
- This is grouping under the make and will indicate the the safest make to invest in/ make that holds a high capacity of passengers

```
In [55]: top_uninjured_make = df.groupby('Make')['Total.Uninjured'].sum().sort_values(ascending=True)
top_uninjured_make
```

```
Out[55]: Make
CESSNA      21023.0
PIPER       10170.0
BEECH       2403.0
BELL        1606.0
ROBINSON     961.0
MOONEY       911.0
BELLANCA     791.0
GRUMMAN      766.0
HUGHES       740.0
MAULE        663.0
Name: Total.Uninjured, dtype: float64
```

```
In [56]: #graph to visualise

plt.figure(figsize = (12,6))
sns.barplot(x = top_uninjured_make.values, y = top_uninjured_make.index)
plt.title('Top 10 Total Uninjured by Aircraft Make')
plt.xlabel('Total Uninjured')
plt.ylabel('Aircraft Make')
plt.show()
```



From the results above, the Cessna brand is the safest make to invest as a company for an aircraft. This is because compared to the rest, it has the largest difference in the number of uninjured people after an accident.

### Broad Phase of Flight against Total Uninjured

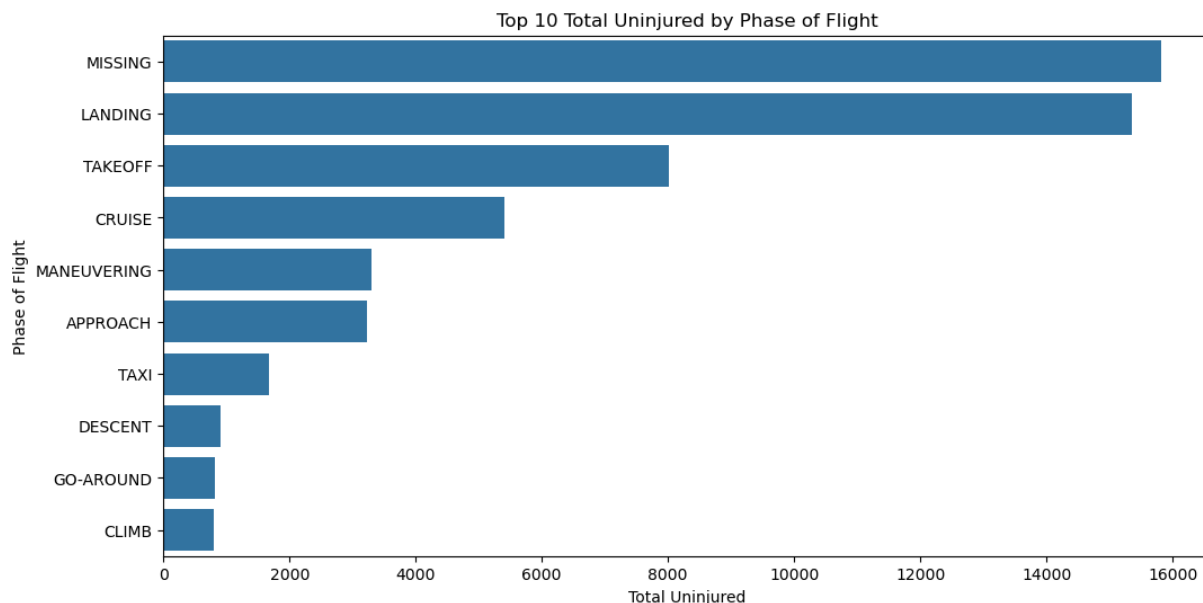
- This is will help draw a conclusion on what the company should focus on (aviation team). This is whether when landing, cruising, etc

```
In [57]: uninjured_by_phase = df.groupby('Broad.phase.of.flight')['Total.Uninjured'].sum().s
uninjured_by_phase
```

```
Out[57]: Broad.phase.of.flight
MISSING      15818.0
LANDING      15362.0
TAKEOFF       8016.0
CRUISE        5402.0
MANEUVERING   3305.0
APPROACH      3230.0
TAXI          1668.0
DESCENT        901.0
GO-AROUND     819.0
CLIMB         795.0
Name: Total.Uninjured, dtype: float64
```

```
In [58]: #graph for visualising

plt.figure(figsize = (12,6))
sns.barplot(x = uninjured_by_phase.values, y = uninjured_by_phase.index)
plt.title('Top 10 Total Uninjured by Phase of Flight')
plt.xlabel('Total Uninjured')
plt.ylabel('Phase of Flight')
plt.show()
```



According to research, there are multiple phases of flight that can be present in aviation. These include the landing, takeoff, cruise and etc. In the case of the results above, we can assume that the results for "MISSING" is Standing phase of flight because it would mean that it is stationary (the aircraft), hence would have the least uninjured individuals out of all the other phases. Out of the 10 featured, the riskiest phases are descent, go around and climb. Visualising and recognising this can help the company pilots be aware.


## Weather Condition Impact on Total Uninjured

- This is to indicate which is the safest weather condition, or rather which weather condition holds the least risk

```
In [59]: uninjured_by_weather = df.groupby('Weather.Condition')['Total.Uninjured'].sum().sort_values()
uninjured_by_weather
```

```
Out[59]: Weather.Condition
VMC      53784.0
MISSING   1076.0
IMC       860.0
UNK       236.0
Name: Total.Uninjured, dtype: float64
```

```
In [60]: #plotting for visuals
plt.figure(figsize = (12,6))
sns.barplot(x = uninjured_by_weather.values, y = uninjured_by_weather.index)
plt.title('Total Uninjured by Weather Condition')
plt.xlabel('Total Uninjured')
plt.ylabel('Weather Condition')
plt.show()
```



The safest weather to travel in based on the results above is VMC which means there is clear weather and good visibility for the pilots. This means that the company would take note of when it is best to have flights taking place and when to avoid flying which would be UNK, meaning unknown i.e, weather is not recorded and unknown.

## Engine Type

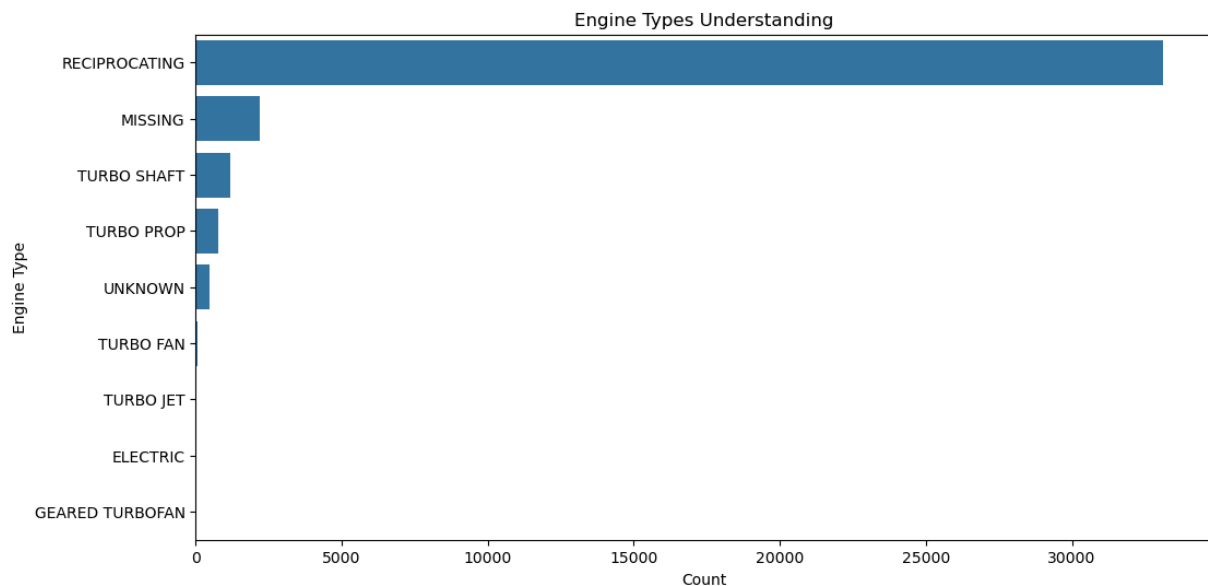
- Understanding the engines involved in accidents

```
In [61]: engine_type_counts = df['Engine.Type'].value_counts()
engine_type_counts
```

```
Out[61]: Engine.Type
RECIPROCATING    33124
MISSING          2188
TURBO SHAFT      1181
TURBO PROP        766
UNKNOWN          454
TURBO FAN         57
TURBO JET         31
ELECTRIC           3
GEARED TURBOFAN    1
Name: count, dtype: int64
```

```
In [62]: #visuals
plt.figure(figsize = (12,6))
sns.barplot(x = engine_type_counts.values, y = engine_type_counts.index)
plt.title('Engine Types Understanding')
plt.xlabel('Count')
plt.ylabel('Engine Type')
plt.show()
```





The results indicate that Reciprocating engine type is one to be avoided as it has the highest count in the number of accidents it is involved in. The safest options for the company to invest in is; turbo fan, turbo jet, electric and geared turbofan

## Flight purpose

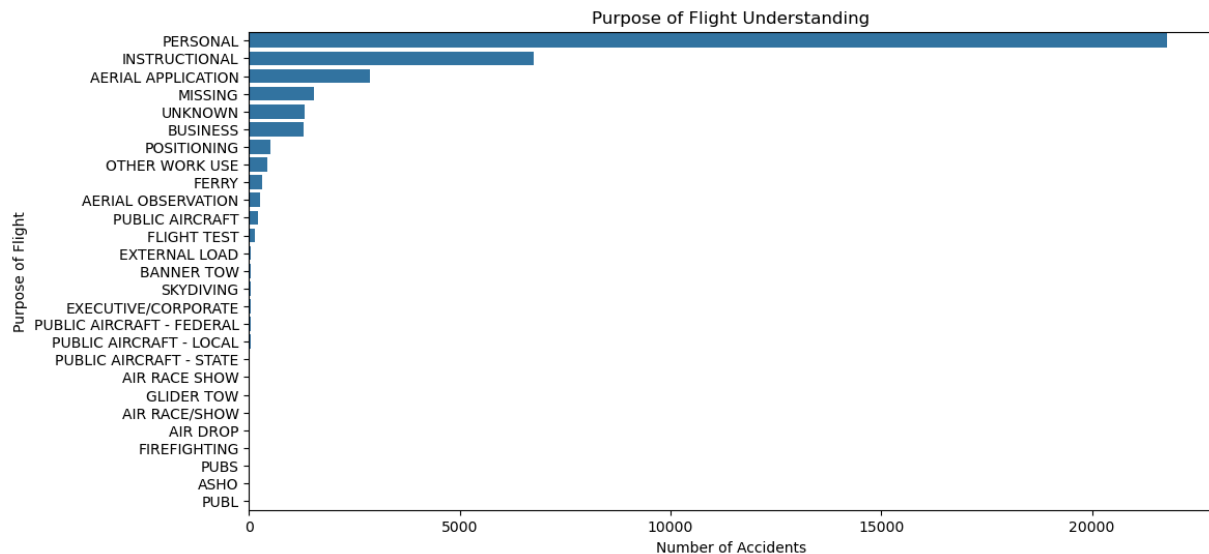
- This will help draw a conclusion on the best purpose to invest in as a company venturing into the aviation sector.

```
In [63]: purpose_counts = df['Purpose.of.flight'].value_counts()  
purpose_counts
```

```
Out[63]: Purpose.of.flight
PERSONAL 21781
INSTRUCTIONAL 6752
AERIAL APPLICATION 2860
MISSING 1552
UNKNOWN 1318
BUSINESS 1299
POSITIONING 519
OTHER WORK USE 443
FERRY 302
AERIAL OBSERVATION 256
PUBLIC AIRCRAFT 224
FLIGHT TEST 153
EXTERNAL LOAD 49
BANNER TOW 47
SKYDIVING 45
EXECUTIVE/CORPORATE 38
PUBLIC AIRCRAFT - FEDERAL 38
PUBLIC AIRCRAFT - LOCAL 31
PUBLIC AIRCRAFT - STATE 30
AIR RACE SHOW 27
GLIDER TOW 18
AIR RACE/SHOW 9
AIR DROP 5
FIREFIGHTING 5
PUBS 2
ASHO 1
PUBL 1
Name: count, dtype: int64
```

```
In [64]: # a visual representation of the data above
```

```
plt.figure(figsize = (12,6))
sns.barplot(x = purpose_counts.values, y = purpose_counts.index)
plt.title('Purpose of Flight Understanding')
plt.xlabel('Number of Accidents')
plt.ylabel('Purpose of Flight')
plt.show()
```



From the results above, it is risky for the company to invest in personal/privatised flights. This is because it has the highest count of accidents compared to all the other purposes.

# Conclusion and Recommendations

## Conclusion

In the analysis above, we explored aviation accident data from 1962 to 2023 to identify insights for aircraft purchase and operational safety recommendations. The data was collected from Kaggle and presented by NTSB. From the cleaning and analysis we have been able to form visuals to help form recommendations for the company.

The dataset was cleaned by:

- Handling missing values strategically (mode or 'Missing')
- Removing or capping outliers using IQR
- Narrowing to top makes and models for focused insights
- Understanding the engine types, weather, and purpose of the flight

Key findings include:

- The frequency of accidents over the years has decreased
- Certain makes and models have higher accident frequencies but also higher uninjured counts, which suggests relative safety in accidents.
- Weather conditions and flight phases significantly impact safety outcomes.
- The type of engine matters based on the number of accidents as well as the purpose on the frequency of accidents that take place

## Recommendations

1. Aircraft Selection: Note and prioritise the makes, engine types and models with high uninjured passenger counts, this indicates that the aircrafts have robust safety features.
2. Operational Focus: Strengthen pilot training and protocols for high-risk phases such as approach and landing as well as handling the unpredictable weather and how to overcome them.
3. Purpose of flight: It's important to invest in flights that have less accident frequency such as the public/commercial flights to avoid any business risk on the company